

MAIN FORECASTING

```
In [4]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.tsa.stattools import adfuller
```

```
In [7]: # Load the dataset and strip column names of extra spaces
data = pd.read_excel(r"C:\Users\KARTHIKEYA\Desktop\Mini Project\Datasets\
data.columns = data.columns.str.strip() # Remove leading/trailing spaces

# Display the first few rows and check column names
print("Columns in dataset:", data.columns)
print(data.head())

# Ensure 'Year' is in datetime format and set as index
if 'Year' in data.columns:
    data['Year'] = pd.to_datetime(data['Year'], format='%Y', errors='coer
    data.set_index('Year', inplace=True)
else:
    print("Error: 'Year' column not found!")

# Identify correct column name for coal consumption
coal_col = None
for col in data.columns:
    if "coal consumption" in col.lower(): # Case-insensitive match
        coal_col = col
        break

if coal_col:
    time_series = data[coal_col]
    print("Coal Consumption Data Extracted Successfully")
else:
    print("Error: Coal Consumption column not found! Available columns:",
```


Columns in dataset: Index(['Entity', 'Year', 'Coal consumption - TWh'], dt
ype='object')

	Entity	Year	Coal consumption - TWh
0	Africa	1965	323.49615
1	Africa	1966	323.12220
2	Africa	1967	330.29156
3	Africa	1968	343.51290
4	Africa	1969	346.64288


Coal Consumption Data Extracted Successfully

```
In [49]: #SCATTER PLOT SHOWING CONSUMPTION OF COAL FROM 1965 - 2023

plt.figure(figsize=(12, 7))
plt.scatter(time_series.index, time_series, color='steelblue', alpha=0.7)
plt.title("Coal Consumption Over Time (Scatter Plot)")
plt.xlabel("Year")
plt.ylabel("Coal Consumption")
plt.grid(True, linestyle="--", alpha=0.5)
plt.show()
```



```
In [54]: #TIME SERIES PLOT(LINE GRAPH) SHOWING CONSUMPTION OF COAL FROM 1965 - 2023
# Visualize the data
plt.figure(figsize=(12, 7))
plt.plot(time_series, label="Coal Consumption")
plt.title("Coal Consumption Over Time")
plt.xlabel("Year")
plt.ylabel("Coal Consumption")
plt.legend()
plt.show()
```



```
In [50]: #COUNTRY WISE CONSUMPTION OF PAST 8 YEARS

import matplotlib.pyplot as plt

# Filter for years 2015-2023
start_year, end_year = 2015, 2023
```

```
filtered_data = data[(data["Year"] >= start_year) & (data["Year"] <= end_

# Get unique countries
countries = filtered_data["Entity"].unique()


# Plot multiple countries with different colors
plt.figure(figsize=(12, 10))

for country in countries:
    country_data = filtered_data[filtered_data["Entity"] == country]
    plt.plot(country_data["Year"], country_data["Coal consumption - TWh"])

# Scatter plot with country labels
for _, row in filtered_data.iterrows():
    plt.text(row["Year"], row["Coal consumption - TWh"], row["Entity"], f

# Improve visualization
plt.title("Coal Consumption Over Time (2015-2023)")
plt.xlabel("Year")
plt.ylabel("Coal Consumption (TWh)")
plt.legend(loc="upper left", bbox_to_anchor=(1, 1))
plt.grid(True, linestyle="--", alpha=0.5)

# Show plot
plt.show()
```



```
In [118... # Check for stationarity using ADF (Augmented Dickey-Fuller) test
adf_test = adfuller(time_series)
print("ADF Test Results:")
print(f"p-value: {adf_test[1]}")

if adfuller(time_series)[1] <= 0.05:
    print("\033[1mTHE DATA IS STATIONARY.\033[0m")
else:
    print("\033[1mTHE DATA IS NOT STATIONARY, APPLYING DIFFERENCING...\033[0m")


# Apply differencing if necessary
time_series_diff = time_series.diff().dropna()

# Plot differenced data
plt.figure(figsize=(13, 7))
plt.plot(time_series_diff, label="Differenced Coal Consumption")
plt.title("Differenced Coal Consumption")
plt.xlabel("Year")
plt.ylabel("Differenced Consumption")
plt.legend()
plt.show()
```

ADF Test Results:

p-value: 0.0

THE DATA IS STATIONARY.



```
In [67]: import pandas as pd

# Load the dataset
df = pd.read_excel(r"C:\Users\KARTHIKEYA\Desktop\Mini Project\Datasets\CO")

# Print the first few rows to check column names
print(df.head())

# Ensure the column names match your dataset
df.rename(columns=lambda x: x.strip(), inplace=True) # Strip any extra spaces

# Assign time series data
india_data = df # Assigning to a variable for consistency
time_series = india_data["Coal consumption - TWh"]
```

	Entity	Year	Coal consumption – TWh
0	Africa	1965	323.49615
1	Africa	1966	323.12220
2	Africa	1967	330.29156
3	Africa	1968	343.51290
4	Africa	1969	346.64288

```
In [69]: # TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP
# import pandas as pd

# # Load the dataset
# file_path = r"C:\Users\KARTHIKEYA\Desktop\Mini Project\Datasets\COAL Co
# df = pd.read_excel(file_path)

# # Clean column names (strip spaces)
# df.columns = df.columns.str.strip()

# # Ensure the correct column name is used
# column_name = "Coal consumption – TWh" # Adjust if needed
# if column_name not in df.columns:
#     print(f"Column '{column_name}' not found. Available columns: {df.co
#     exit()

# # Assign time series data
# time_series = df[column_name].dropna()
```

```
In [80]: # Fit ARIMA model
p, d, q = 1, 2, 1 # Example parameters (can be tuned)
arma_model = SARIMAX(time_series, order=(p, d, q), seasonal_order=(0, 0,
arma_results = arma_model.fit()

# Summary of ARIMA model
print(arma_results.summary())

# Forecast next 5 years using ARIMA
forecast_arma = arma_results.get_forecast(steps=5)
forecast_values_arma = forecast_arma.predicted_mean
confidence_intervals_arma = forecast_arma.conf_int()
```

SARIMAX Results

```

=====
=====
Dep. Variable:      Coal consumption - TWh   No. Observations:
5521
Model:              SARIMAX(1, 2, 1)         Log Likelihood      -4
6480.771
Date:              Fri, 14 Feb 2025         AIC                  9
2967.543
Time:              23:44:06                 BIC                  9
2987.390
Sample:            0                       HQIC               9
2974.464
                                - 5521
Covariance Type:    opg
=====
=====

```

	coef	std err	z	P> z	[0.025	0.
975]						

ar.L1	-0.0076	0.035	-0.220	0.826	-0.076	
0.060						
ma.L1	-0.9999	0.005	-213.604	0.000	-1.009	-
0.991						
sigma2	1.209e+06	6481.975	186.467	0.000	1.2e+06	1.22
e+06						

```

=====
=====
Ljung-Box (L1) (Q):      0.00   Jarque-Bera (JB):      149
065679.55
Prob(Q):                 1.00   Prob(JB):
0.00
Heteroskedasticity (H):  0.37   Skew:
-24.47
Prob(H) (two-sided):     0.00   Kurtosis:
806.64
=====
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```

In [87]: # Fit SARIMA model
P, D, Q, s = 1, 1, 1, 2 # Example seasonal parameters (can be tuned)
sarima_model = SARIMAX(time_series, order=(p, d, q), seasonal_order=(P, D, Q, s))
sarima_results = sarima_model.fit()

# Summary of SARIMA model
print(sarima_results.summary())

# Forecast next 5 years using SARIMA
forecast_sarima = sarima_results.get_forecast(steps=5)
forecast_values_sarima = forecast_sarima.predicted_mean
confidence_intervals_sarima = forecast_sarima.conf_int()

```

SARIMAX Results

```

=====
=====
Dep. Variable:          Coal consumption - TWh    No. Observations:
5521
Model:                SARIMAX(1, 2, 1)x(1, 1, 1, 2)    Log Likelihood
-46488.942
Date:                  Fri, 14 Feb 2025    AIC
92987.885
Time:                  23:46:25    BIC
93020.963
Sample:                0    HQIC
92999.421
                        - 5521
Covariance Type:      opg
=====
=====

```

	coef	std err	z	P> z	[0.025	0.
975]						

ar.L1	-0.0074	0.033	-0.223	0.824	-0.072	
0.057						
ma.L1	-1.0000	0.833	-1.201	0.230	-2.632	
0.632						
ar.S.L2	0.0118	0.028	0.420	0.675	-0.043	
0.067						
ma.S.L2	-1.0000	0.833	-1.201	0.230	-2.632	
0.632						
sigma2	1.104e+06	2.53e-08	4.36e+13	0.000	1.1e+06	1.1
e+06						
=====						
=====						
Ljung-Box (L1) (Q):		0.00	Jarque-Bera (JB):		151	
179498.25						
Prob(Q):		0.98	Prob(JB):			
0.00						
Heteroskedasticity (H):		0.37	Skew:			
-24.62						
Prob(H) (two-sided):		0.00	Kurtosis:			
812.47						
=====						
=====						

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 4.68e+29. Standard errors may be unstable.

```

In [128... # Print forecasts (ALL COUNTRIES)
print("\nARIMA Forecast for Next 5 Years:")
print(forecast_values_arma)
print("\nSARIMA Forecast for Next 5 Years:")
print(forecast_values_sarima)

```


ARIMA Forecast for Next 5 Years:

```
5521    45569.594535
5522    45579.532620
5523    45589.430468
5524    45599.328623
5525    45609.226775
```

Name: predicted_mean, dtype: float64

SARIMA Forecast for Next 5 Years:

```
2024-01-01    45947.859348
2025-01-01    46691.167951
2026-01-01    47111.347328
2027-01-01    47886.204829
2028-01-01    48319.895327
```

Freq: YS-JAN, Name: predicted_mean, dtype: float64

```
In [111... # TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP TEMP 1
# import matplotlib.pyplot as plt

# # Generate future years for plotting
# future_years = pd.date_range(start=time_series.index[-1] + pd.DateOffset(
# plt.figure(figsize=(10, 6))
# plt.plot(time_series, label="Historical Data")
# plt.plot(future_years, forecast_values_arima, label="ARIMA Forecast", l
# plt.plot(future_years, forecast_values_sarima, label="SARIMA Forecast",
# plt.fill_between(future_years, confidence_intervals_arima.iloc[:, 0], c
# plt.fill_between(future_years, confidence_intervals_sarima.iloc[:, 0],
# plt.title(f"{entity} Coal Consumption Forecast")
# plt.xlabel("Year")
# plt.ylabel("Coal Consumption (TWh)")
# plt.legend()
# plt.show()
```

```
In [134... # Assuming you have loaded your data and filtered for a SPECIFIC COUNTRY
entity = "Algeria"

# Your existing code for loading, preprocessing, and fitting models...

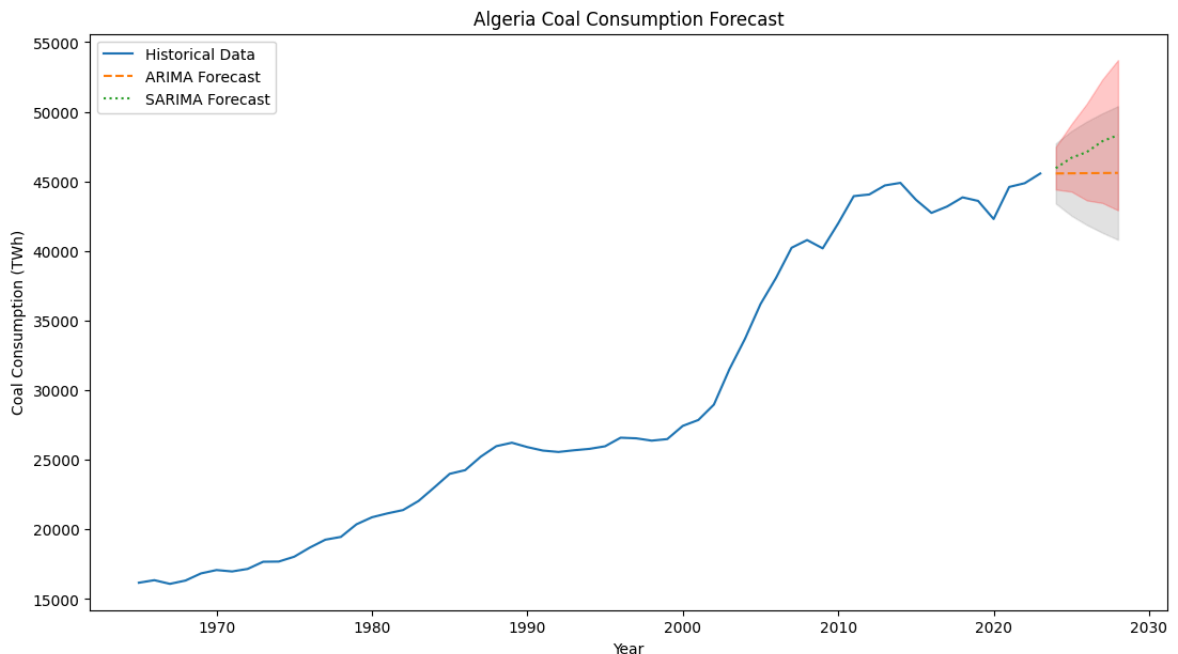
# Forecast next 5 years using ARIMA
forecast_arima = arima_results.get_forecast(steps=5)
forecast_values_arima = forecast_arima.predicted_mean
confidence_intervals_arima = forecast_arima.conf_int()

# Forecast next 5 years using SARIMA
forecast_sarima = sarima_results.get_forecast(steps=5)
forecast_values_sarima = forecast_sarima.predicted_mean
confidence_intervals_sarima = forecast_sarima.conf_int()

# Generate future years for plotting
future_years = pd.date_range(
    start=time_series.index[-1] + pd.offsets.YearBegin(),
    periods=5,
    freq='YS'
)

# Plot forecasts
plt.figure(figsize=(13, 7))
plt.plot(time_series, label="Historical Data")
plt.plot(future_years, forecast_values_arima, label="ARIMA Forecast", lin
```

```
plt.plot(future_years, forecast_values_sarima, label="SARIMA Forecast", l
plt.fill_between(future_years, confidence_intervals_arima.iloc[:, 0], con
plt.fill_between(future_years, confidence_intervals_sarima.iloc[:, 0], co
plt.title(f"{entity} Coal Consumption Forecast")
plt.xlabel("Year")
plt.ylabel("Coal Consumption (TWh)")
plt.legend()
plt.show()
```



WHOLE CODE FOR ANY RANDOM COUNTRY REPORT GENERATION

```
In [139... import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.tsa.stattools import adfuller

# Load the dataset
# file_path = "3c681988-d6d8-4378-ba4c-2486974e9c91_COAL Consumption.xlsx"
# sheet_name = "coal-consumption-by-country-ter"
# data = pd.read_excel(file_path, sheet_name=sheet_name)

data = pd.read_excel(r"C:\Users\KARTHIKEYA\Desktop\Mini Project\Datasets\

# Display the first few rows of the dataset
print(data.head())

# Function to preprocess and analyze data for a given entity
def analyze_entity(entity):
    # Filter data for the specific entity
    entity_data = data[data['Entity'] == entity].copy()

    # Ensure 'Year' is datetime and set as index
    entity_data['Year'] = pd.to_datetime(entity_data['Year'], format='%Y')
    entity_data.set_index('Year', inplace=True)

    # Extract the coal consumption column
    time_series = entity_data['Coal consumption - TWh']
```

```

# Check if there is enough data to proceed
if len(time_series) < 20:
    print(f"Not enough data for {entity}. Please select another count")
    return

# Visualize the data
plt.figure(figsize=(10, 6))
plt.plot(time_series, label=f"{entity} Coal Consumption")
plt.title(f"{entity} Coal Consumption Over Time")
plt.xlabel("Year")
plt.ylabel("Coal Consumption (TWh)")
plt.legend()
plt.show()

# Check for stationarity using ADF test
adf_test = adfuller(time_series)
print("ADF Test Results:")
print(f"p-value: {adf_test[1]}")
if adf_test[1] <= 0.05:
    print("The data is stationary.")
else:
    print("The data is non-stationary. Applying differencing...")

# Apply differencing if necessary
time_series_diff = time_series.diff().dropna()

# Plot differenced data
plt.figure(figsize=(10, 6))
plt.plot(time_series_diff, label="Differenced Coal Consumption")
plt.title(f"Differenced {entity} Coal Consumption")
plt.xlabel("Year")
plt.ylabel("Differenced Consumption")
plt.legend()
plt.show()

# Fit ARIMA model
p, d, q = 1, 1, 1 # Example parameters (can be tuned)
arima_model = SARIMAX(time_series, order=(p, d, q), seasonal_order=(0
arima_results = arima_model.fit()

# Summary of ARIMA model
print(arima_results.summary())

# Forecast next 5 years using ARIMA
forecast_arima = arima_results.get_forecast(steps=5)
forecast_values_arima = forecast_arima.predicted_mean
confidence_intervals_arima = forecast_arima.conf_int()

# Fit SARIMA model
P, D, Q, s = 1, 1, 1, 2 # Example seasonal parameters (can be tuned)
sarima_model = SARIMAX(time_series, order=(p, d, q), seasonal_order=(
sarima_results = sarima_model.fit()

# Summary of SARIMA model
print(sarima_results.summary())

# Forecast next 5 years using SARIMA
forecast_sarima = sarima_results.get_forecast(steps=5)
forecast_values_sarima = forecast_sarima.predicted_mean
confidence_intervals_sarima = forecast_sarima.conf_int()

```

```


# Generate future years for plotting
future_years = pd.date_range(
    start=time_series.index[-1] + pd.offsets.YearBegin(),
    periods=5,
    freq='YS'
)

# Plot forecasts
plt.figure(figsize=(13, 7))
plt.plot(time_series, label="Historical Data")
plt.plot(future_years, forecast_values_arima, label="ARIMA Forecast",
plt.plot(future_years, forecast_values_sarima, label="SARIMA Forecast")
plt.fill_between(future_years, confidence_intervals_arima.iloc[:, 0],
plt.fill_between(future_years, confidence_intervals_sarima.iloc[:, 0]
plt.title(f"{entity} Coal Consumption Forecast")
plt.xlabel("Year")
plt.ylabel("Coal Consumption (TWh)")
plt.legend()
plt.show()

# Example usage: Analyze data for India
entity = "India"
analyze_entity(entity)

```


	Entity	Year	Coal consumption - TWh
0	Africa	1965	323.49615
1	Africa	1966	323.12220
2	Africa	1967	330.29156
3	Africa	1968	343.51290
4	Africa	1969	346.64288



ADF Test Results:

p-value: 1.0

The data is non-stationary. Applying differencing...



```
C:\Users\KARTHIKEYA\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency YS-JAN will be used.
```

```
self._init_dates(dates, freq)
```

```
C:\Users\KARTHIKEYA\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency YS-JAN will be used.
```

```
self._init_dates(dates, freq)
```

```
C:\Users\KARTHIKEYA\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:966: UserWarning: Non-stationary starting autoregressive parameters found. Using zeros as starting parameters.
```

```
warn('Non-stationary starting autoregressive parameters')
```

SARIMAX Results

```

=====
=====
Dep. Variable:      Coal consumption - TWh    No. Observations:
59
Model:              SARIMAX(1, 1, 1)          Log Likelihood
-362.543
Date:              Sat, 15 Feb 2025          AIC
731.087
Time:              00:37:01                  BIC
737.268
Sample:            01-01-1965                HQIC
733.495
                  - 01-01-2023
Covariance Type:   opg
=====
=====

```

	coef	std err	z	P> z	[0.025	0.
975]						
ar.L1	0.9927	0.017	60.113	0.000	0.960	
1.025						
ma.L1	-0.8489	0.068	-12.493	0.000	-0.982	-
0.716						
sigma2	1.527e+04	1603.095	9.523	0.000	1.21e+04	1.84
e+04						

```

=====
=====
Ljung-Box (L1) (Q):      1.46    Jarque-Bera (JB):
171.07
Prob(Q):                 0.23    Prob(JB):
0.00
Heteroskedasticity (H):  59.17    Skew:
0.91
Prob(H) (two-sided):     0.00    Kurtosis:
11.21
=====
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

C:\Users\KARTHIKEYA\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency YS-JAN will be used.

self._init_dates(dates, freq)

C:\Users\KARTHIKEYA\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency YS-JAN will be used.

self._init_dates(dates, freq)

SARIMAX Results

```

=====
=====
Dep. Variable:          Coal consumption - TWh    No. Observations:
59
Model:                SARIMAX(1, 1, 1)x(1, 1, 1, 2)    Log Likelihood
-349.252
Date:                  Sat, 15 Feb 2025    AIC
708.505
Time:                  00:37:02    BIC
718.631
Sample:                01-01-1965    HQIC
712.431
                        - 01-01-2023
Covariance Type:      opg
=====
=====

```

	coef	std err	z	P> z	[0.025	0.
975]						

ar.L1	-0.9907	12.344	-0.080	0.936	-25.185	2
3.204						
ma.L1	0.9977	25.259	0.039	0.968	-48.510	5
0.505						
ar.S.L2	-0.8222	0.185	-4.443	0.000	-1.185	-
0.460						
ma.S.L2	0.3034	0.289	1.051	0.293	-0.262	
0.869						
sigma2	1.493e+04	1.95e+05	0.077	0.939	-3.67e+05	3.97
e+05						

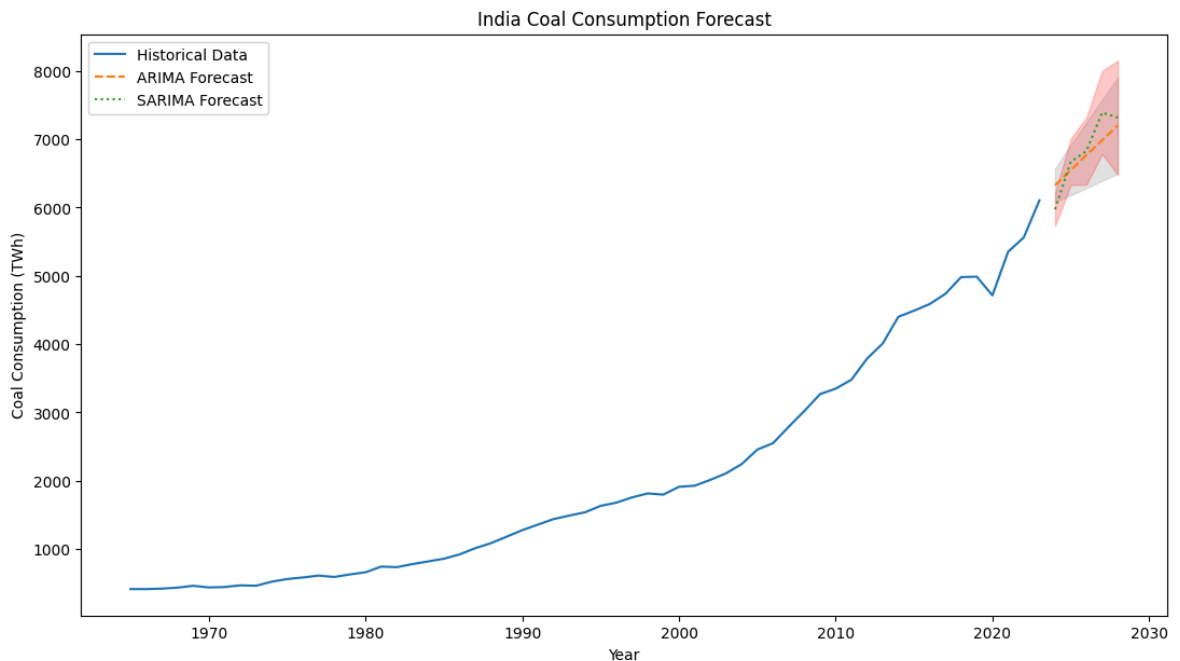
```

=====
=====
Ljung-Box (L1) (Q):      0.07    Jarque-Bera (JB):
121.41
Prob(Q):                0.80    Prob(JB):
0.00
Heteroskedasticity (H): 33.48    Skew:
0.72
Prob(H) (two-sided):    0.00    Kurtosis:
10.07
=====
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).



WHOLE CODE FOR ANY RANDOM COUNTRY REPORT GENERATION WITH EVALUATION METRICS (MAE, RMSE, MAPE)

```
In [144... import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.tsa.stattools import adfuller

# Load the dataset
data = pd.read_excel(r"C:\Users\KARTHIKEYA\Desktop\Mini Project\Datasets\

# Display the first few rows of the dataset
print(data.head())

# Function to preprocess and analyze data for a given entity
def analyze_entity(entity):
    # Filter data for the specific entity
    entity_data = data[data['Entity'] == entity].copy()

    # Ensure 'Year' is datetime and set as index
    entity_data['Year'] = pd.to_datetime(entity_data['Year'], format='%Y')
    entity_data.set_index('Year', inplace=True)

    # Extract the coal consumption column
    time_series = entity_data['Coal consumption - TWh']

    # Check if there is enough data to proceed
    if len(time_series) < 20:
        print(f"Not enough data for {entity}. Please select another count")
        return

    # Visualize the data
    plt.figure(figsize=(10, 6))
    plt.plot(time_series, label=f"{entity} Coal Consumption")
    plt.title(f"{entity} Coal Consumption Over Time")
    plt.xlabel("Year")
```



```

plt.ylabel("Coal Consumption (TWh)")
plt.legend()
plt.show()

# Check for stationarity using ADF test
adf_test = adfuller(time_series)
print("ADF Test Results:")
print(f"p-value: {adf_test[1]}")
if adf_test[1] <= 0.05:
    print("The data is stationary.")
else:
    print("The data is non-stationary. Applying differencing...")

# Apply differencing if necessary
time_series_diff = time_series.diff().dropna()

# Fit ARIMA model
p, d, q = 1, 1, 1 # Example parameters (can be tuned)
arima_model = SARIMAX(time_series, order=(p, d, q), seasonal_order=(0
arima_results = arima_model.fit()

# Summary of ARIMA model
print(arima_results.summary())

# Forecast next 5 years using ARIMA
forecast_arima = arima_results.get_forecast(steps=5)
forecast_values_arima = forecast_arima.predicted_mean
confidence_intervals_arima = forecast_arima.conf_int()

# Fit SARIMA model
P, D, Q, s = 1, 1, 1, 2 # Example seasonal parameters (can be tuned)
sarima_model = SARIMAX(time_series, order=(p, d, q), seasonal_order=(
sarima_results = sarima_model.fit()

# Summary of SARIMA model
print(sarima_results.summary())

# Forecast next 5 years using SARIMA
forecast_sarima = sarima_results.get_forecast(steps=5)
forecast_values_sarima = forecast_sarima.predicted_mean
confidence_intervals_sarima = forecast_sarima.conf_int()

# Calculate performance metrics
actual_values = time_series.values
forecast_values_arima = forecast_values_arima.values
forecast_values_sarima = forecast_values_sarima.values

# Mean Absolute Error (MAE) for ARIMA and SARIMA
mae_arima = np.mean(np.abs(actual_values[-len(forecast_values_arima):
mae_sarima = np.mean(np.abs(actual_values[-len(forecast_values_sarima

# Root Mean Squared Error (RMSE) for ARIMA and SARIMA
rmse_arima = np.sqrt(np.mean((actual_values[-len(forecast_values_arim
rmse_sarima = np.sqrt(np.mean((actual_values[-len(forecast_values_sar

# Mean Absolute Percentage Error (MAPE) for ARIMA and SARIMA
mape_arima = np.mean(np.abs((actual_values[-len(forecast_values_arima
mape_sarima = np.mean(np.abs((actual_values[-len(forecast_values_sari

# Calculate performance metrics as percentages

```

```

total_actual_sum = np.sum(actual_values)
mae_arima_percentage = (mae_arima / total_actual_sum) * 100
rmse_arima_percentage = (rmse_arima / total_actual_sum) * 100
mape_arima_percentage = mape_arima / 100

mae_sarima_percentage = (mae_sarima / total_actual_sum) * 100
rmse_sarima_percentage = (rmse_sarima / total_actual_sum) * 100
mape_sarima_percentage = mape_sarima / 100

# Print performance metrics
print("\nPerformance Metrics for ARIMA:")
print(f"MAE (Percentage): {mae_arima_percentage:.2f}%")
print(f"RMSE (Percentage): {rmse_arima_percentage:.2f}%")
print(f"MAPE (Percentage): {mape_arima_percentage:.2f}%")

print("\nPerformance Metrics for SARIMA:")
print(f"MAE (Percentage): {mae_sarima_percentage:.2f}%")
print(f"RMSE (Percentage): {rmse_sarima_percentage:.2f}%")
print(f"MAPE (Percentage): {mape_sarima_percentage:.2f}%")


# Generate future years for plotting
future_years = pd.date_range(
    start=time_series.index[-1] + pd.offsets.YearBegin(),
    periods=5,
    freq='YS'
)

# Plot forecasts
plt.figure(figsize=(13, 7))
plt.plot(time_series, label="Historical Data")
plt.plot(future_years, forecast_values_arima, label="ARIMA Forecast",
plt.plot(future_years, forecast_values_sarima, label="SARIMA Forecast",
plt.fill_between(future_years, confidence_intervals_arima.iloc[:, 0],
plt.fill_between(future_years, confidence_intervals_sarima.iloc[:, 0],
plt.title(f"{entity} Coal Consumption Forecast")
plt.xlabel("Year")
plt.ylabel("Coal Consumption (TWh)")
plt.legend()
plt.show()

# Example usage: Analyze data for India
entity = "India"
analyze_entity(entity)

```

	Entity	Year	Coal consumption - TWh
0	Africa	1965	323.49615
1	Africa	1966	323.12220
2	Africa	1967	330.29156
3	Africa	1968	343.51290
4	Africa	1969	346.64288



ADF Test Results:

p-value: 1.0

The data is non-stationary. Applying differencing...

```
C:\Users\KARTHIKEYA\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency YS-JAN will be used.
```

```
self._init_dates(dates, freq)
```

```
C:\Users\KARTHIKEYA\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency YS-JAN will be used.
```

```
self._init_dates(dates, freq)
```

```
C:\Users\KARTHIKEYA\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\statespace\sarimax.py:966: UserWarning: Non-stationary starting autoregressive parameters found. Using zeros as starting parameters.
```

```
warn('Non-stationary starting autoregressive parameters')
```

SARIMAX Results

```

=====
=====
Dep. Variable:      Coal consumption - TWh    No. Observations:
59
Model:              SARIMAX(1, 1, 1)          Log Likelihood
-362.543
Date:              Sat, 15 Feb 2025          AIC
731.087
Time:              00:42:43                  BIC
737.268
Sample:            01-01-1965                HQIC
733.495
                  - 01-01-2023
Covariance Type:    opg
=====
=====

```

	coef	std err	z	P> z	[0.025	0.
975]						

ar.L1	0.9927	0.017	60.113	0.000	0.960	
1.025						
ma.L1	-0.8489	0.068	-12.493	0.000	-0.982	-
0.716						
sigma2	1.527e+04	1603.095	9.523	0.000	1.21e+04	1.84
e+04						

```

=====
=====
Ljung-Box (L1) (Q):      1.46    Jarque-Bera (JB):
171.07
Prob(Q):                 0.23    Prob(JB):
0.00
Heteroskedasticity (H):  59.17    Skew:
0.91
Prob(H) (two-sided):     0.00    Kurtosis:
11.21
=====
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

C:\Users\KARTHIKEYA\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency YS-JAN will be used.

self._init_dates(dates, freq)

C:\Users\KARTHIKEYA\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency YS-JAN will be used.

self._init_dates(dates, freq)

SARIMAX Results

```

=====
Dep. Variable:          Coal consumption - TWh    No. Observations:
59
Model:                 SARIMAX(1, 1, 1)x(1, 1, 1, 2)    Log Likelihood
-349.252
Date:                  Sat, 15 Feb 2025    AIC
708.505
Time:                  00:42:43    BIC
718.631
Sample:                01-01-1965    HQIC
712.431
                        - 01-01-2023
Covariance Type:      opg
=====

```

```

=====
              coef      std err          z      P>|z|      [0.025      0.
975]
-----
ar.L1          -0.9907      12.344      -0.080      0.936      -25.185      2
3.204
ma.L1           0.9977      25.259       0.039      0.968      -48.510      5
0.505
ar.S.L2        -0.8222       0.185     -4.443      0.000       -1.185      -
0.460
ma.S.L2         0.3034       0.289       1.051      0.293       -0.262
0.869
sigma2         1.493e+04    1.95e+05     0.077      0.939     -3.67e+05    3.97
e+05
=====

```

```

=====
Ljung-Box (L1) (Q):          0.07    Jarque-Bera (JB):
121.41
Prob(Q):                    0.80    Prob(JB):
0.00
Heteroskedasticity (H):      33.48    Skew:
0.72
Prob(H) (two-sided):         0.00    Kurtosis:
10.07
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

Performance Metrics for ARIMA:

MAE (Percentage): 1.18%
 RMSE (Percentage): 1.19%
 MAPE (Percentage): 0.27%

Performance Metrics for SARIMA:

MAE (Percentage): 1.23%
 RMSE (Percentage): 1.27%
 MAPE (Percentage): 0.28%

